

HW1

Ravenclaws

2023-09-24

Q1 - drop NA from me1 and me2

dependencies and env. settings

```
library(reticulate)
use_condaenv('jiaxin')

library(dplyr)
library(tidyr)
library(ggplot2)
library(GGally)
library(e1071)
library(caret)

library(parallel)
numCores = detectCores()

library(doParallel)
registerDoParallel(numCores-1)

library(keras)
library(tensorflow)
tf <- import('tensorflow')
seed = 0590
set.seed(seed)
py_set_seed(seed, disable_hash_randomization = TRUE)

set.a.seed = function() {
  set_random_seed(
    seed,
    disable_gpu = TRUE
  )
}
```

preprocessing

```
## read in data

data = read.csv('sdss-all-c.csv', header = F)
colname = c('id', 'class', 'brightness', 'size', 'texture', 'me1', 'me2')
colnames(data) <- colname
data$class[data$class == 3] <- 0 # galaxy
data$class[data$class == 6] <- 1 # star
data$me1 = as.numeric(data$me1)
data$me2 = as.numeric(data$me2)

## drop NA observations
data_complete7 = data %>% drop_na()
data_complete5 = data[, 1:5]

## standardized
data_complete7[, 3:7] = scale(data_complete7[, 3:7])
data_complete5[, 3:5] = scale(data_complete5[, 3:5])

## pack datasets
packing = function(data) {

  set.seed()
  ## randomly partition the dataset (.75 training)
  num = dim(data)[1]
  training_idx = sample(1:num, num*0.75, replace = F)
  data.train <- data[training_idx, ]
  data.test <- data[-training_idx, ]

  ## gathering data
  train.inputs <- data.train[, 3:dim(data.train)[2]]
  train.targets = as.matrix(data.train$class)
  test.inputs <- data.test[, 3:dim(data.test)[2]]
  test.targets = as.matrix(data.test$class)

  dataset = list()
  dataset$train = list(input = train.inputs, targets = train.targets)
  dataset$test = list(input = test.inputs, targets = test.targets)
  dataset$trainsvm = data.train
  dataset$testsvm = data.test

  return(dataset)
}

## check classification
table(data_complete7$class)
```

```
##  
##    0    1  
## 1178  287
```

```
## get dataset  
dataset = packing(data_complete7)  
str(dataset)
```

```
## List of 4  
## $ train :List of 2  
##   ..$ input :'data.frame': 1098 obs. of  5 variables:  
##     ..$ brightness: num [1:1098] -0.188 0.104 0.999 0.704 0.509 ...  
##     ..$ size      : num [1:1098] 1.6645 -0.0572 0.2818 -0.7612 -0.3632 ...  
##     ..$ texture   : num [1:1098] -0.118 -0.121 -0.122 -0.12 -0.12 ...  
##     ..$ me1       : num [1:1098] 0.126 -0.19 -0.47 -0.341 0.18 ...  
##     ..$ me2       : num [1:1098] 0.506 -0.016 0.377 -0.103 0.234 ...  
##     ..$ targets: num [1:1098, 1] 0 0 0 0 0 0 0 0 0 0 ...  
## $ test  :List of 2  
##   ..$ input :'data.frame': 367 obs. of  5 variables:  
##     ..$ brightness: num [1:367] -2.526 -3.001 -1.565 -0.718 -2.354 ...  
##     ..$ size      : num [1:367] -0.584 -0.582 -0.568 0.124 -0.582 ...  
##     ..$ texture   : num [1:367] 0.38 2.391 -0.1 -0.118 0.392 ...  
##     ..$ me1       : num [1:367] -0.006889 0.000529 -0.007486 -0.016787 -0.022757 ...  
##     ..$ me2       : num [1:367] -0.0155 -0.0197 -0.0215 0.1036 -0.0266 ...  
##     ..$ targets: num [1:367, 1] 1 1 1 0 1 1 1 1 1 1 ...  
## $ trainsvm:'data.frame': 1098 obs. of  7 variables:  
##   ..$ id      : chr [1:1098] "77-446" "77-739" "76-1239" "75-592" ...  
##   ..$ class   : num [1:1098] 0 0 0 0 0 0 0 0 0 0 ...  
##   ..$ brightness: num [1:1098] -0.188 0.104 0.999 0.704 0.509 ...  
##   ..$ size     : num [1:1098] 1.6645 -0.0572 0.2818 -0.7612 -0.3632 ...  
##   ..$ texture  : num [1:1098] -0.118 -0.121 -0.122 -0.12 -0.12 ...  
##   ..$ me1     : num [1:1098] 0.126 -0.19 -0.47 -0.341 0.18 ...  
##   ..$ me2     : num [1:1098] 0.506 -0.016 0.377 -0.103 0.234 ...  
## $ testvsm :'data.frame': 367 obs. of  7 variables:  
##   ..$ id      : chr [1:367] "75-32" "75-38" "75-42" "75-49" ...  
##   ..$ class   : num [1:367] 1 1 1 0 1 1 1 1 1 1 ...  
##   ..$ brightness: num [1:367] -2.526 -3.001 -1.565 -0.718 -2.354 ...  
##   ..$ size     : num [1:367] -0.584 -0.582 -0.568 0.124 -0.582 ...  
##   ..$ texture  : num [1:367] 0.38 2.391 -0.1 -0.118 0.392 ...  
##   ..$ me1     : num [1:367] -0.006889 0.000529 -0.007486 -0.016787 -0.022757 ...  
##   ..$ me2     : num [1:367] -0.0155 -0.0197 -0.0215 0.1036 -0.0266 ...
```

linear classification by tensorflow

wrapped functions

```

sq_loss = function(targets, pred) {
  tf$reduce_mean(tf$square(tf$subtract(targets, pred)))
}

# model
model = function(inputs) {tf$matmul(inputs, w) + b}

## run step
run = function(t_input, targets, lr) {
  with(tf$GradientTape() %as% tape, {
    pred <- model(t_input)
    loss <- sq_loss(targets, pred)
  })
  update <- tape$gradient(loss, list(w=w, b=b))

  w$assign_sub(update$w * lr)
  b$assign_sub(update$b * lr)
  loss
}

## training
model_train = function(t_input, targets, lr, printwidth=5) {
  step = 0
  thres = 1e-6
  curr = 0

  repeat {

    step <- step + 1
    t_loss <- run(t_input, targets, lr)
    loss <- as.numeric(sprintf("%.3f\n", t_loss))

    if (step %% printwidth == 0) cat(sprintf("Loss at step %s: %.6f\n", step, t_loss))
    updates <- abs(loss - curr)

    if (step > 0 & updates <= thres) {
      cat(sprintf('Finish after %s steps.\n', step))
      break
    }
    curr <- loss
  }
}

## evaluation
eval = function(targets, pred) {
  mse = sum(targets - pred)^2 / dim(targets)[1]
  pred = ifelse(pred >= 0.5, 1, 0)
}

```

```
acc = sum(targets == pred)
cat(sprintf('Accuracy: %.6f(%s/%s)  MSE: %.6f\n\n', acc/dim(targets)[1], acc, dim(targets)[1], mse))
targets = as.factor(targets)
pred = as.factor(round(pred))
print(confusionMatrix(pred, targets)$table)
}
```

training result (Qa.i)

```
lr_list = c(0.01, 0.05, 0.1, 0.15, 0.2)

input_dim = dim(dataset$train$input)[2]
output_dim = dim(dataset$train$targets)[2]

for (lr in lr_list){
  cat('Learning rate: ', lr, '\n')

  # initialize hidden layer parameters
  set.seed()
  w <- tf$Variable(tf$random$uniform(shape(input_dim, output_dim)))
  b <- tf$Variable(tf$zeros(shape(output_dim)))
  train.t_input = as_tensor(dataset$train$input, dtype='float32')

  # train
  model_train(train.t_input, dataset$train$targets, lr)

  # eval
  test.t_input = as_tensor(dataset$test$input, dtype='float32')
  prediction = as.vector(model(test.t_input))
  eval(dataset$test$targets, prediction)
  cat('-----\n')
}
```

```
## Learning rate: 0.01
## Loss at step 5: 0.924639
## Loss at step 10: 0.747607
## Loss at step 15: 0.609703
## Loss at step 20: 0.501660
## Loss at step 25: 0.416504
## Loss at step 30: 0.348972
## Loss at step 35: 0.295077
## Loss at step 40: 0.251789
## Loss at step 45: 0.216799
## Loss at step 50: 0.188336
## Loss at step 55: 0.165038
## Loss at step 60: 0.145852
## Loss at step 65: 0.129960
## Loss at step 70: 0.116724
## Loss at step 75: 0.105641
## Loss at step 80: 0.096316
## Loss at step 85: 0.088432
## Loss at step 90: 0.081739
## Loss at step 95: 0.076034
## Loss at step 100: 0.071153
## Finish after 103 steps.
## Accuracy: 0.888283(326/367) MSE: 0.127740
##
## Reference
## Prediction 0 1
##          0 292 39
##          1  2  34
## -----
## Learning rate: 0.05
## Loss at step 5: 0.464817
## Loss at step 10: 0.199318
## Loss at step 15: 0.108626
## Loss at step 20: 0.072033
## Loss at step 25: 0.055428
## Loss at step 30: 0.047318
## Finish after 33 steps.
## Accuracy: 0.959128(352/367) MSE: 0.001036
##
## Reference
## Prediction 0 1
##          0 294 15
##          1  0  58
## -----
## Learning rate: 0.1
## Loss at step 5: 0.214873
## Loss at step 10: 0.073169
## Loss at step 15: 0.047319
## Loss at step 20: 0.040912
## Finish after 22 steps.
## Accuracy: 0.972752(357/367) MSE: 0.018473
##
```

```
##           Reference
## Prediction 0   1
##          0 294 10
##          1   0  63
## -----
## Learning rate: 0.15
## Loss at step 5: 0.116835
## Loss at step 10: 0.047304
## Loss at step 15: 0.039687
## Finish after 15 steps.
## Accuracy: 0.975477(358/367) MSE: 0.021835
##
##           Reference
## Prediction 0   1
##          0 294  9
##          1   0  64
## -----
## Learning rate: 0.2
## Loss at step 5: 0.075587
## Loss at step 10: 0.040731
## Finish after 13 steps.
## Accuracy: 0.980926(360/367) MSE: 0.027474
##
##           Reference
## Prediction 0   1
##          0 294   7
##          1   0  66
## -----
```

According to the evaluation result, the best linear classification model achieves an 98.1% accuracy on the test dataset with learning rate = 0.2 and updates only for 13 steps. According to the change in confusion matrix, increasing learning rate from 0.01 to 0.2 enables the model to recall some false negative predictions.

classification by svm (Qa.ii, Qa.iii)

```
formula = class~brightness+size+texture+me1+me2

set.seed(seed)
kernels = c('linear', 'radial', 'polynomial', 'sigmoid')
costs = c(0.001, 0.01, 0.1, 1, 5, 10, 100)

# parallel tuning process
params = expand.grid(kernel = kernels, cost = costs)
results = list()
foreach(i = 1:nrow(params), .packages = 'e1071') %do% {
  current_cost = params$cost[i]
  current_kernel = params$kernel[i]
  curResult <- tune(method = svm,
                     train.x = formula,
                     data = dataset$trainsvm,
                     kernel = as.character(current_kernel),
                     cost = current_cost)
  results[[i]] <- c(current_kernel, current_cost, curResult$performances$error)
}
```

```
## [[1]]
## [1] 1.000000 0.001000 0.112261
##
## [[2]]
## [1] 2.0000000 0.0010000 0.1670796
##
## [[3]]
## [1] 3.0000000 0.0010000 0.1600346
##
## [[4]]
## [1] 4.0000000 0.0010000 0.1658615
##
## [[5]]
## [1] 1.00000000 0.01000000 0.04292841
##
## [[6]]
## [1] 2.00000000 0.01000000 0.07274986
##
## [[7]]
## [1] 3.0000000 0.0100000 0.1088974
##
## [[8]]
## [1] 4.0000000 0.0100000 0.0922463
##
## [[9]]
## [1] 1.00000000 0.10000000 0.04047963
##
## [[10]]
## [1] 2.00000000 0.10000000 0.01899529
##
## [[11]]
## [1] 3.0000000 0.1000000 0.1124262
##
## [[12]]
## [1] 4.0000000 0.1000000 0.2156642
##
## [[13]]
## [1] 1.00000000 1.00000000 0.04031662
##
## [[14]]
## [1] 2.000000000 1.000000000 0.008682732
##
## [[15]]
## [1] 3.0000000 1.0000000 0.2578538
##
## [[16]]
## [1] 4.00000 1.00000 34.78067
##
## [[17]]
## [1] 1.00000000 5.00000000 0.04022732
##
## [[18]]
```

```

## [1] 2.000000000 5.000000000 0.006386856
##
## [[19]]
## [1] 3.0000000 5.0000000 0.3200779
##
## [[20]]
## [1] 4.000 5.000 1032.743
##
## [[21]]
## [1] 1.0000000 10.0000000 0.04054598
##
## [[22]]
## [1] 2.000000000 10.000000000 0.005846379
##
## [[23]]
## [1] 3.000000 10.000000 1.919723
##
## [[24]]
## [1] 4.000 10.000 5020.248
##
## [[25]]
## [1] 1.000000000 100.000000000 0.04057071
##
## [[26]]
## [1] 2.000000e+00 1.000000e+02 5.673254e-03
##
## [[27]]
## [1] 3.000000 100.000000 2.112995
##
## [[28]]
## [1] 4.0 100.0 428140.4

```

```

tuning_result = data.frame(Kernel = levels(params$kernel)[results[[i]]][1],
                           Cost = results[[i]][2],
                           Error = results[[i]][3])

# packing results
for (i in 2:nrow(params)) {
  tuning_result = rbind(tuning_result,
                        c(levels(params$kernel)[results[[i]]][1],
                          results[[i]][2],
                          results[[i]][3]))
}

tuning_result$Kernel = as.factor(tuning_result$Kernel)
tuning_result$Cost = as.factor(tuning_result$Cost)
tuning_result>Error = as.numeric(tuning_result>Error)

tuning_result[order(tuning_result>Error, decreasing = F), ]

```

	Kernel <fct>	Cost <fct>	Error <dbl>
26	radial	100	5.673254e-03
22	radial	10	5.846379e-03
18	radial	5	6.386856e-03
14	radial	1	8.682732e-03
10	radial	0.1	1.899529e-02
17	linear	5	4.022732e-02
13	linear	1	4.031662e-02
9	linear	0.1	4.047963e-02
21	linear	10	4.054598e-02
25	linear	100	4.057071e-02

1-10 of 28 rows

Previous 1 2 3 Next

Use 10-fold cross validation to calculate the average error of each model (under different parameter combination). According to the tuning result, radial kernel separates the data best. And the best cost parameter is 100.

```
best.svm = svm(formula, dataset$trainsvm, kernel = 'radial', cost = 100)
summary(best.svm)
```

```
##
## Call:
## svm(formula = formula, data = dataset$trainsvm, kernel = "radial",
##       cost = 100)
##
##
## Parameters:
##   SVM-Type:  eps-regression
##   SVM-Kernel: radial
##     cost: 100
##     gamma: 0.2
##   epsilon: 0.1
##
##
## Number of Support Vectors:  286
```

```
pred = predict(best.svm, dataset$testvsm)
eval(dataset$testvsm$class, pred)
```

```
##           Reference
## Prediction  0   1
##           0 294   0
##           1   0  73
```

Using the best model to evaluate model performance on test dataset. By checking the confusion matrix, the model achieves 100% accuracy on the test dataset.

Q1 - drop column me1 and me2 (Qb)

```
dataset = packing(data_complete5)
```

classification by network

```
lr_list = c(0.01, 0.05, 0.1, 0.15, 0.2)
input_dim = dim(dataset$train$input)[2]
output_dim = dim(dataset$train$targets)[2]

for (lr in lr_list){
  cat('Learning rate: ', lr, '\n')

  # initialize hidden layer parameters
  set.seed()
  w <- tf$Variable(tf$random$uniform(shape(input_dim, output_dim)))
  b <- tf$Variable(tf$zeros(shape(output_dim)))
  train.t_input = as_tensor(dataset$train$input, dtype='float32')

  # train
  model_train(train.t_input, dataset$train$targets, lr)

  # eval
  test.t_input = as_tensor(dataset$test$input, dtype='float32')
  prediction = as.vector(model(test.t_input))
  eval(dataset$test$targets, prediction)
  cat('-----\n')
}
```

```
## Learning rate: 0.01
## Loss at step 5: 0.847400
## Loss at step 10: 0.675107
## Loss at step 15: 0.543668
## Loss at step 20: 0.442800
## Loss at step 25: 0.364888
## Loss at step 30: 0.304282
## Loss at step 35: 0.256782
## Loss at step 40: 0.219258
## Loss at step 45: 0.189369
## Loss at step 50: 0.165362
## Loss at step 55: 0.145914
## Loss at step 60: 0.130027
## Loss at step 65: 0.116941
## Loss at step 70: 0.106078
## Loss at step 75: 0.096993
## Loss at step 80: 0.089340
## Loss at step 85: 0.082853
## Loss at step 90: 0.077321
## Loss at step 95: 0.072579
## Finish after 95 steps.
## Accuracy: 0.862069(325/377) MSE: 0.469356
##
##          Reference
## Prediction 0 1
##           0 299 52
##           1 0 26
## -----
## Learning rate: 0.05
## Loss at step 5: 0.407417
## Loss at step 10: 0.174044
## Loss at step 15: 0.099457
## Loss at step 20: 0.069361
## Loss at step 25: 0.055135
## Loss at step 30: 0.047836
## Finish after 34 steps.
## Accuracy: 0.949602(358/377) MSE: 0.004470
##
##          Reference
## Prediction 0 1
##           0 299 19
##           1 0 59
## -----
## Learning rate: 0.1
## Loss at step 5: 0.186121
## Loss at step 10: 0.070509
## Loss at step 15: 0.047907
## Loss at step 20: 0.041772
## Finish after 20 steps.
## Accuracy: 0.973475(367/377) MSE: 0.000710
##
##          Reference
```

```
## Prediction 0 1
##          0 299 10
##          1  0  68
## -----
## Learning rate: 0.15
## Loss at step 5: 0.106171
## Loss at step 10: 0.047969
## Loss at step 15: 0.040557
## Finish after 15 steps.
## Accuracy: 0.976127(368/377) MSE: 0.004903
##
##           Reference
## Prediction 0 1
##          0 299 9
##          1  0  69
## -----
## Learning rate: 0.2
## Loss at step 5: 0.073120
## Loss at step 10: 0.041610
## Finish after 13 steps.
## Accuracy: 0.978780(369/377) MSE: 0.009468
##
##           Reference
## Prediction 0 1
##          0 299 8
##          1  0  70
## -----
```

According to the evaluation result, the best linear classification model achieves a 97.9% accuracy on the test dataset with learning rate = 0.2 and updates only for 13 steps. According to the change in confusion matrix, increasing learning rate from 0.01 to 0.2 enables the model to recall some false negative predictions. Comparing with Qa's data cleaning strategy, utilizing the information in variable me1 and me2 in linear classification can increase model's accuracy by around 1%.

classification by svm

```
formula = class~brightness+size+texture

set.seed(seed)
kernels = c('linear', 'radial', 'polynomial', 'sigmoid')
costs = c(0.001, 0.01, 0.1, 1, 5, 10, 100)

# parallel tuning process
params = expand.grid(kernel = kernels, cost = costs)
results = list()
foreach(i = 1:nrow(params), .packages = 'e1071') %do% {
  current_cost = params$cost[i]
  current_kernel = params$kernel[i]
  curResult <- tune(method = svm,
                     train.x = formula,
                     data = dataset$trainsvm,
                     kernel = as.character(current_kernel),
                     cost = current_cost)
  results[[i]] <- c(current_kernel, current_cost, curResult$performances$error)
}
```

```
## [[1]]
## [1] 1.0000000 0.0010000 0.1126463
##
## [[2]]
## [1] 2.0000000 0.0010000 0.155732
##
## [[3]]
## [1] 3.0000000 0.0010000 0.118142
##
## [[4]]
## [1] 4.0000000 0.0010000 0.1553271
##
## [[5]]
## [1] 1.00000000 0.01000000 0.04373643
##
## [[6]]
## [1] 2.0000000 0.0100000 0.0504816
##
## [[7]]
## [1] 3.00000000 0.01000000 0.08488147
##
## [[8]]
## [1] 4.00000000 0.01000000 0.09005643
##
## [[9]]
## [1] 1.00000000 0.10000000 0.04143525
##
## [[10]]
## [1] 2.00000000 0.10000000 0.01309013
##
## [[11]]
## [1] 3.0000000 0.1000000 0.1730617
##
## [[12]]
## [1] 4.0000000 0.1000000 0.7721243
##
## [[13]]
## [1] 1.00000000 1.00000000 0.04114575
##
## [[14]]
## [1] 2.000000000 1.000000000 0.005843049
##
## [[15]]
## [1] 3.0000000 1.0000000 0.6051436
##
## [[16]]
## [1] 4.00000 1.00000 56.58255
##
## [[17]]
## [1] 1.00000000 5.00000000 0.04159578
##
## [[18]]
```

```

## [1] 2.000000000 5.000000000 0.004194124
##
## [[19]]
## [1] 3.0000000 5.0000000 0.6953375
##
## [[20]]
## [1] 4.000 5.000 1419.663
##
## [[21]]
## [1] 1.0000000 10.0000000 0.04121756
##
## [[22]]
## [1] 2.000000000 10.000000000 0.004331692
##
## [[23]]
## [1] 3.0000000 10.0000000 0.5852891
##
## [[24]]
## [1] 4.000 10.000 5379.161
##
## [[25]]
## [1] 1.0000000 100.0000000 0.04113837
##
## [[26]]
## [1] 2.000000e+00 1.000000e+02 2.735516e-03
##
## [[27]]
## [1] 3.0000000 100.0000000 0.6034993
##
## [[28]]
## [1] 4.0 100.0 755142.7

```

```

tuning_result = data.frame(Kernel = levels(params$kernel)[results[[i]]][1],
                           Cost = results[[i]][2],
                           Error = results[[i]][3])

# packing results
for (i in 2:nrow(params)) {
  tuning_result = rbind(tuning_result,
                        c(levels(params$kernel)[results[[i]]][1],
                          results[[i]][2],
                          results[[i]][3]))
}

tuning_result$Kernel = as.factor(tuning_result$Kernel)
tuning_result$Cost = as.factor(tuning_result$Cost)
tuning_result>Error = as.numeric(tuning_result>Error)

tuning_result[order(tuning_result>Error, decreasing = F), ]

```

	Kernel <fct>	Cost <fct>	Error <dbl>
26	radial	100	2.735516e-03
18	radial	5	4.194124e-03
22	radial	10	4.331692e-03
14	radial	1	5.843049e-03
10	radial	0.1	1.309013e-02
25	linear	100	4.113837e-02
13	linear	1	4.114575e-02
21	linear	10	4.121756e-02
9	linear	0.1	4.143525e-02
17	linear	5	4.159578e-02

1-10 of 28 rows

Previous 1 2 3 Next

```
best.svm = svm(formula, dataset$trainsvm, kernel = 'radial', cost = 100)
summary(best.svm)
```

```
##
## Call:
## svm(formula = formula, data = dataset$trainsvm, kernel = "radial",
##       cost = 100)
##
##
## Parameters:
##   SVM-Type:  eps-regression
##   SVM-Kernel: radial
##     cost: 100
##     gamma: 0.3333333
##     epsilon: 0.1
##
##
## Number of Support Vectors: 184
```

```
pred = predict(best.svm, dataset$testvsm)
eval(dataset$testvsm$class, pred)
```

```
##           Reference
## Prediction 0 1
##           0 299 0
##           1 0 78
```

Use 10-fold cross validation to calculate the average error of each model (under different parameter combination). According to the tuning result, radial kernel separates the data best. And the best cost parameter is 100. The best model also achieves 100% accuracy on test dataset. Two different data cleaning strategies do not influence the model performance of svm. Since under each strategy, svm can separate the data perfectly.

Q2

preprocessing

```

x = read.table('ziptrain.dat')
y = read.table('zipdigit.dat')

## reformat input (2000 rows, 256 columns)
data <- cbind(y, x)

## create 75 25 split by class
set.seed(seed)
split <- createDataPartition(y = data[,1], p = 0.75) |> unlist()
train <- data[split,]
test <- data[-split,]
rm(split, data)

## set up data sets, scale training data set
train <- list(input = train[,2:ncol(train)], targets = train[,1])
train[['input']] <- train[['input']] |> sapply(as.numeric) |> as.matrix()
train[['targets']] <- train[['targets']] |> as.matrix()
scale <- train[['input']] |> preProcess(method = c('center', 'scale'))
train[['input']] <- predict(scale, train[['input']])

## optionally scale test data based on training data
test <- list(input = test[,2:ncol(test)], targets = test[,1])
test[['input']] <- test[['input']] |> sapply(as.numeric) |> as.matrix()
test[['targets']] <- test[['targets']] |> as.matrix()
test[['input']] <- predict(scale, test[['input']]) |> as.matrix()

## make sure distribution of digits is similar between training and test sets
par(mfrow = c(1, 2))
hist(train$targets)
hist(test$targets)

```



PCA(Qa.i)

```
# wrapped function
summary_image_by = function(func) {
  summary_img = array(dim = c(10, 256))
  for (i in 0:9) {
    summary_img[i+1, ] <- apply(x[which(y == i), ], 2, func)
  }
  summary_img
}

mean_img = summary_image_by(mean)

# remove mean effect
mean_effect = array(dim = c(nrow(x), 256))
for (i in 0:9) {
  nrows = length(which(y == i))
  mean_effect[which(y == i), ] <- matrix(rep(mean_img[i+1, ], nrows), nrow = nrows, byrow = T)
}

cleaned_zip_img = x - mean_effect

zip.pca = prcomp(cleaned_zip_img)

curr_var = 0
Divider = TRUE
total_var = sum((zip.pca$sdev)^2)
compress_idx = 1

cat("Cumulative Importance:\n")
```

```
## Cumulative Importance:
```

```
for (i in 1:length(zip.pca$sdev)) {
  curr_var = curr_var + zip.pca$sdev[i]^2
  cat(sprintf("%s - %.2f\n", i, curr_var/total_var*100))
  if (curr_var / total_var > 0.8 & Divider) {
    cat('-----\n')
    Divider = FALSE
    compress_idx = i
    break
  }
}
```

```
## 1 - 10.51
## 2 - 18.39
## 3 - 24.13
## 4 - 28.89
## 5 - 33.33
## 6 - 37.15
## 7 - 40.64
## 8 - 43.29
## 9 - 45.73
## 10 - 48.03
## 11 - 50.20
## 12 - 52.10
## 13 - 53.93
## 14 - 55.68
## 15 - 57.31
## 16 - 58.85
## 17 - 60.36
## 18 - 61.72
## 19 - 63.02
## 20 - 64.31
## 21 - 65.49
## 22 - 66.63
## 23 - 67.73
## 24 - 68.76
## 25 - 69.68
## 26 - 70.58
## 27 - 71.44
## 28 - 72.30
## 29 - 73.14
## 30 - 73.94
## 31 - 74.71
## 32 - 75.45
## 33 - 76.16
## 34 - 76.86
## 35 - 77.50
## 36 - 78.12
## 37 - 78.71
## 38 - 79.30
## 39 - 79.86
## 40 - 80.40
## -----
```

After removing the effect of digit-specific means and applying PCA, the first 46 components can explain at least 80% of the total variance in the training dataset.

low-rank representation (Qa.ii)

```
# wrapped function
replace_d = function(d, k) {
  d_k = diag(0, nrow = length(d))
  d_k <- rep(0, length(d))
  d_k[1:k] <- d[1:k]
  diag(d_k)
}

renewed_img_of = function(d_k) {
  decomp_img = svd_res$u %*% d_k %*% t(svd_res$v)
  renewed_img = decomp_img + mean_effect
  dim(renewed_img) <- c(nrow(x), 16, 16)
  renewed_img
}
```

```
# svd
svd_res = svd(cleaned_zip_img)

# display
d_k = replace_d(svd_res$d, 40)
new_img = renewed_img_of(d_k)

# raw form
par(mar = rep(0.05, 4), mfrow = c(40, 50))
raw_img = as.matrix(x)
dim(raw_img) <- c(nrow(x), 16, 16)
apply(raw_img, 1, function(x) {image(x[, 16:1], axes = F, col = gray(31:0/31))})
```

654734310170111774801487487374186741377454274137740
 6320966408782090220812083382208144898967619708046
 80030809038012290665920919117109080791384435168544
 68448640239868935680226841027102271092704808721132
 73222702285412278702710260276081711017764462911903
 11810316117553650688000149570001601986019868009685
 44940864812121268021301321380279000068917949199453
 85718014655320106446090260559617776192060321618666
 010260256090260521605466473618-1609026043560561820
 60902618016052261761604731186090261701605220927609
 02618216191360902619004626180146992971028070297102
 82342821228658287392821728373009328365830328601281
 0282652826284328232811232721222116270224012345230
 609300230088300833024030815302743080930030230247811
 11213721038037130851685708911180005200112003759402
 97418213002211021090214601720019706021380197001840
 13640260101830910498072006509006300090700934406910
 70806600918654006040075800230066009010092300912006
 97075423557188600180202048280602132971040185023853
 45891073815246206010746912232048838666149700697405
 98002909840849550482322698480840720850065370692576
 57196546854468009680036811485027857188522711893190
 11997719714191197141972602066222463164631466547390
 01900457006270801700940296405054101724079413384254
 14170141531409144301421014311414724194230142501425
 01411391470321332120221268212021262140421701121228
 30366303263035603493030263034430327199661980319713
 19519861980419713672177748670627241725037251272203
 720322174059815138134370753813838117372029607370838
 0501989919968197201980403197111971117717404176011
 70131737317011175731736017837175421706617604300177
 63757316915734011705617311757317306170217268174071
 73331701117055197201972600119804197031971198501970
 79700198039851989101961906102277110019726001019805
 19897989914201802171101972600100540105345054010540
 10564108870288203710062260603806790087060913506103
 06087060870409604963492004751536314014106049049301
 42009720197207971029721018642193201933518049181011
 5918901187041910119132191011901192551910419124191
 41190131400319130191011903110110319010107190151901

```
## NULL
```

```
# low-rank form
par(mar = rep(0.05, 4), mfrow = c(40, 50))
apply(new_img, 1, function(x) {image(x[, 16:1], axes = F, col = gray(31:0/31))})
```

```

65473031012011174801497487314186741377454274137740
6320866806782090208120833082208144898967619708046
80030809038012290665920919117109080791384435168544
68448640239848935680226841027102271092704808727132
73220710228512279702710260276082711017764462911903
11810316117553650688000149570001601986011868007685
44940864312121230213023211880279000068917999197453
85718014635320106646090240559617776182060521617666
0102602560902605216054664736181609026043560561820
609026180160522617616047311826090261701605220724609
02618216191340902619004626180146972971028070292103
823418212286582873708217283720007228365830328601281
028265282628432833281128202211222116270224013345230
609300230088300833024030815302743080930030230247811
11213121038037130851483708911180005200112003759402
9741821300221402109021460172001706213801970018410
136402601018309104780720065090063000907009300912006
7080660091865400604007580023006600901009300912006
97075423567188600180201048280605132971040185023853
45891073815296306010746912232048838666149700671405
98004909840849550482322698480810720850065370692574
591196468544680091688036811485027857118322711873190
11997719714191197141972608066222468164631466547390
019004570062708017200940296405054101724079415924254
17170141531409144301421014311414224194230142501425
01411391476321230120221268218021342140421701121228
30366303263035603493030263034430327199661950319713
19519061780419713672177798670627941725037251222903
72032314059815138134370753813838117372029607370838
05019899199681172019880403197111971117917404176011
7013131757317011175731736017537175421706617609300177
637573169157340117056173117573117300170217268174071
73331701117055177201972600119804197031971198501970
7470011980398811989101961906102277110019716001019805
198979899142018021714019726001100540105345058010590
10564108870288203110062260603806790087060918306103
06081060870409004963492004751536314011166049049301
#20097201972027921029721018642193201933518049181011
59189011870419101191321910119101192551910419124191
#1190131400314130141011903110110319010107190151901

```

```
## NULL
```

By applying PCA to obtain the low-rank representation of the digits, the comparison with raw digit images implies that the PCA can still maintain the shape of each digit, while it may lose some information about the boundary between the digit and its background, reflecting by the vagueness in low-rank representation.

classification performance (Qb.i)

```
# evaluation function for 10 classes
eval10 = function(targets, pred) {
  mse = sum(targets - pred)^2 / dim(targets)[1]
  pred = ifelse(pred >= 9, 9, ifelse(pred <= 0, 0, pred))
  pred = round(pred)
  acc = sum(targets == pred)
  cat(sprintf('Accuracy: %.6f(%s/%s)  MSE: %.6f\n\n', acc/dim(targets)[1], acc, dim(targets)[1], mse))
  targets = as.factor(targets)
  pred = as.factor(round(pred))
  print(confusionMatrix(pred, targets)$table)
}

# packing data
dataset = list()
dataset$train = train
dataset$test = test
```

linear classification

```
lr_list = c(5e-4, 1e-3, 5e-3, 0.01, 0.02)

input_dim = dim(dataset$train$input)[2]
output_dim = dim(dataset$train$targets)[2]

for (lr in lr_list){
  cat('Learning rate: ', lr, '\n')

  # initialize hidden layer parameters
  set.seed()
  w <- tf$Variable(tf$random$uniform(shape(input_dim, output_dim)))
  b <- tf$Variable(tf$zeros(shape(output_dim)))
  train.t_input = as_tensor(dataset$train$input, dtype='float32')

  # train
  model_train(train.t_input, dataset$train$targets, lr, printwidth = 100)

  # eval
  test.t_input = as_tensor(dataset$test$input, dtype='float32')
  prediction = as.vector(model(test.t_input))
  eval10(dataset$test$targets, prediction)
  cat('-----\n')
}
```

```
## Learning rate: 5e-04
## Loss at step 100: 71.352928
## Loss at step 200: 31.162800
## Loss at step 300: 20.819736
## Loss at step 400: 16.445597
## Loss at step 500: 13.808612
## Loss at step 600: 11.920169
## Loss at step 700: 10.468922
## Loss at step 800: 9.318789
## Loss at step 900: 8.392312
## Loss at step 1000: 7.637970
## Loss at step 1100: 7.018681
## Loss at step 1200: 6.506660
## Loss at step 1300: 6.080601
## Loss at step 1400: 5.723929
## Loss at step 1500: 5.423608
## Loss at step 1600: 5.169303
## Loss at step 1700: 4.952754
## Loss at step 1800: 4.767328
## Loss at step 1900: 4.607664
## Loss at step 2000: 4.469415
## Loss at step 2100: 4.349038
## Loss at step 2200: 4.243637
## Finish after 2229 steps.
## Accuracy: 0.290581(145/499)  MSE: 66.860890
##
##          Reference
## Prediction 0 1 2 3 4 5 6 7 8 9
##           0 41 19 15 1 1 0 0 0 0 0
##           1 22 55 3 4 2 0 2 1 2 1
##           2 13 11 7 10 1 1 1 1 1 1
##           3 11 1 12 7 2 2 9 1 4 0
##           4 1 0 9 8 3 5 12 4 3 1
##           5 2 1 5 3 7 4 11 8 8 6
##           6 0 0 3 2 4 6 3 17 10 6
##           7 1 0 1 1 6 2 1 16 7 12
##           8 0 0 1 0 3 1 0 3 3 9
##           9 0 0 0 0 3 0 0 5 1 6
## -----
## Learning rate: 0.001
## Loss at step 100: 31.223835
## Loss at step 200: 16.462408
## Loss at step 300: 11.930202
## Loss at step 400: 9.324646
## Loss at step 500: 7.641434
## Loss at step 600: 6.508751
## Loss at step 700: 5.725228
## Loss at step 800: 5.170140
## Loss at step 900: 4.767887
## Loss at step 1000: 4.469810
## Loss at step 1100: 4.243931
## Loss at step 1200: 4.068899
```

```

## Loss at step 1300: 3. 930245
## Loss at step 1400: 3. 818043
## Finish after 1432 steps.
## Accuracy: 0. 300601(150/499) MSE: 18. 404907
##
##          Reference
## Prediction 0 1 2 3 4 5 6 7 8 9
##          0 36 14 12 1 0 0 0 0 0 0
##          1 24 57 6 2 3 0 2 0 1 0
##          2 16 13 5 11 1 0 1 2 1 2
##          3 9 2 8 10 2 3 6 0 3 0
##          4 3 0 14 5 1 3 15 4 5 1
##          5 2 1 4 4 8 6 11 6 6 4
##          6 0 0 5 2 5 5 4 14 10 5
##          7 1 0 1 1 5 3 0 20 7 13
##          8 0 0 1 0 4 1 0 5 5 11
##          9 0 0 0 0 3 0 0 5 1 6
## -----
## Learning rate: 0.005
## Loss at step 100: 7. 669304
## Loss at step 200: 4. 472988
## Loss at step 300: 3. 726601
## Loss at step 400: 3. 427592
## Loss at step 500: 3. 257137
## Loss at step 600: 3. 142018
## Finish after 612 steps.
## Accuracy: 0. 282565(141/499) MSE: 0. 274156
##
##          Reference
## Prediction 0 1 2 3 4 5 6 7 8 9
##          0 33 9 9 0 1 0 1 0 0 0
##          1 25 51 5 3 1 0 1 0 0 0
##          2 18 23 9 10 2 1 0 0 1 0
##          3 8 2 8 10 2 4 3 2 3 2
##          4 6 1 12 6 1 2 19 1 4 1
##          5 1 1 6 5 6 6 11 10 4 1
##          6 0 0 5 2 6 5 3 12 13 7
##          7 0 0 1 0 5 2 1 17 8 11
##          8 0 0 0 0 6 1 0 8 6 15
##          9 0 0 1 0 2 0 0 6 0 5
## -----
## Learning rate: 0.01
## Loss at step 100: 4. 476960
## Loss at step 200: 3. 428589
## Loss at step 300: 3. 142545
## Loss at step 400: 2. 992941
## Finish after 453 steps.
## Accuracy: 0. 310621(155/499) MSE: 0. 409410
##
##          Reference
## Prediction 0 1 2 3 4 5 6 7 8 9
##          0 34 5 8 0 1 0 2 0 0 0

```

```

##      1 30 55  5  3  1  0  0  0  0  0
##      2 12 23 10 10  2  1  1  0  1  0
##      3  8  2  9 10  2  5  4  2  1  1
##      4  6  1 11  6  3  1 16  1  5  2
##      5  1  1  5  5  4  7 11 11  4  1
##      6  0  0  6  2  6  5  4 10 11  6
##      7  0  0  1  0  5  1  1 18 10 12
##      8  0  0  0  0  5  1  0  6  7 13
##      9  0  0  1  0  3  0  0  8  0  7
## -----
## Learning rate: 0.02
## Loss at step 100: 3.430590
## Loss at step 200: 2.993566
## Loss at step 300: 2.835768
## Finish after 318 steps.
## Accuracy: 0.314629(157/499) MSE: 0.695320
##
##          Reference
## Prediction 0 1 2 3 4 5 6 7 8 9
##      0 35 7 7 0 1 0 2 0 0 0
##      1 29 54 7 4 2 0 0 0 0 0
##      2 13 22 9 8 1 2 1 0 1 0
##      3 8 3 9 12 2 3 5 2 1 1
##      4 4 0 11 7 4 2 16 2 6 2
##      5 2 1 4 3 3 6 10 9 3 2
##      6 0 0 7 2 8 6 5 12 10 5
##      7 0 0 1 0 3 1 0 17 10 12
##      8 0 0 0 0 5 1 0 8 8 13
##      9 0 0 1 0 3 0 0 6 0 7
## -----

```

According to the evaluation result, the best linear classification model achieves a 31.5% accuracy on the test dataset with learning rate = 0.02 and updates for 318 steps (tuning learning rate from 5e-4 to 0.02, since training encountered gradient exploding with use lr =0.05). It implies that the linear classification model performs not well on digit dataset.

SVM

```

colnames(dataset$train$targets) <- c('class')
colnames(dataset$test$targets) <- c('class')
dataset$trainsvm = data.frame(cbind(dataset$train$targets, dataset$train$input))
dataset$testsvm = data.frame(cbind(dataset$test$targets, dataset$test$input))

set.seed(seed)
kernels = c('linear', 'radial', 'polynomial', 'sigmoid')
costs = c(0.001, 0.01, 0.1, 1, 5, 10, 100)

# parallel tuning process
## run each model type with generic cost = 10

results = list()
foreach(i = 1:4, .packages = 'e1071') %do% {
  current_kernel = kernels[i]
  curResult <- tune(method = svm,
                     train.x = class~.,
                     data = dataset$trainsvm,
                     kernel = as.character(current_kernel),
                     cost = 10)
  results[[i]] <- c(current_kernel, 10, curResult$performances$error)
}

```

```

## [[1]]
## [1] "linear"          "10"           "4.56010835336013"
##
## [[2]]
## [1] "radial"          "10"           "0.986574213712546"
##
## [[3]]
## [1] "polynomial"      "10"           "1.35209299278766"
##
## [[4]]
## [1] "sigmoid"         "10"           "8106.70239085516"

```

```
tuning_result = data.frame(Kernel = results[[1]][1],
                           Cost = results[[1]][2],
                           Error = results[[1]][3])

# packing results
for (i in 2:length(kernels)){
  tuning_result = rbind(tuning_result,
                        c(results[[i]][1],
                          results[[i]][2],
                          results[[i]][3]))
}

tuning_result$Kernel = as.factor(tuning_result$Kernel)
tuning_result$Cost = as.factor(tuning_result$Cost)
tuning_result>Error = as.numeric(tuning_result>Error)

tuning_result[order(tuning_result>Error, decreasing = F), ]
```

	Kernel <fct>	Cost <fct>	Error <dbl>
2	radial	10	0.9865742
3	polynomial	10	1.3520930
1	linear	10	4.5601084
4	sigmoid	10	8106.7023909
4 rows			

```
## further tuning
svmfit <- tune(method = svm,
                 train.x = dataset$train$input,
                 train.y = dataset$train$targets,
                 ranges = list(cost = c(0.1, 1, 5, 10, 100),
                               kernel = c('linear', 'radial', 'polynomial', 'sigmoid')),
                 scale = FALSE)

summary(svmfit$best.model)
```

```

## 
## Call:
## best.tune(method = svm, train.x = dataset$train$input, train.y = dataset$train$targets,
##           ranges = list(cost = c(0.1, 1, 5, 10, 100), kernel = c("linear",
##                         "radial", "polynomial", "sigmoid")), scale = FALSE)
##
##
## Parameters:
##   SVM-Type:  eps-regression
##   SVM-Kernel: radial
##     cost: 100
##     gamma: 0.00390625
##   epsilon: 0.1
##
##
## Number of Support Vectors: 1217

```

```

best.svm = svm(class~, dataset$trainsvm, kernel = 'radial', cost = 100)
pred = predict(best.svm, dataset$testsvm)
eval10(dataset$test$targets, pred)

```

```

## Accuracy: 0.609218 (304/499)  MSE: 0.493060
## 
##          Reference
## Prediction 0 1 2 3 4 5 6 7 8 9
##           0 67 0 0 0 0 0 0 0 0 0
##           1 18 82 2 0 0 0 0 0 0 0
##           2 4 4 24 3 0 0 0 0 0 0
##           3 2 1 18 18 4 0 1 0 0 0
##           4 0 0 7 11 16 7 5 1 2 0
##           5 0 0 4 3 9 5 11 6 2 1
##           6 0 0 1 1 1 7 19 6 5 3
##           7 0 0 0 0 2 2 3 37 9 2
##           8 0 0 0 0 0 0 0 6 15 15
##           9 0 0 0 0 0 0 0 0 6 21

```

By grid-search and 10-fold cross validation, the best svm classifier uses radial kernel with cost parameter equal to 100. The best model has 60.9% accuracy on test dataset.