

Q1

We first try using the original images to train the model, but we cannot get an ideal result.

| Best lr | filter | train acc | val acc |
|---------|------------------------------|-----------|---------|
| 0.05 | c(8, 16) | 0.7533 | 0.6333 |
| 0.001 | c(16, 32) | 1.0000 | 0.6200 |
| 0.001 | c(8, 16, 32) | 0.9600 | 0.6667 |
| 0.01 | c(16, 32, 64) | 0.6667 | 0.7133 |
| 0.001 | c(8, 16, 32, 64) | 0.9467 | 0.7600 |
| 0.001 | c(16, 32, 64, 128) | 0.9533 | 0.7067 |
| 0.001 | c(16, 32, 64, 128, 256) | 0.7133 | 0.7000 |
| 0.001 | c(32, 64, 128, 256, 512) | 0.9667 | 0.7667 |
| 0.001 | c(16, 32, 64, 128, 256, 512) | 0.6400 | 0.7333 |

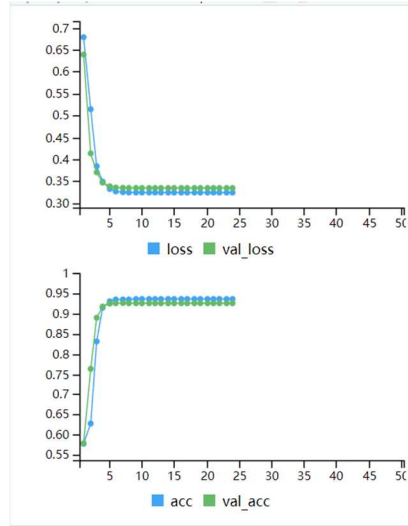
None of the model structures and hyper-parameter settings can help the model get an accuracy over 80% on the validation set. To improve our model's classification performance, we focus on the preprocessing process and achieve unattainable scores compared to using the original dataset alone. We extract and sample small images from each image, utilizing these image snippets that potentially contain individual letters or at least a few letters. Specifically, we randomly sample some small pieces with an acceptance rule that the small piece has less than 10% white pixels to make sure that the sub picture is not whitespace. This sampling approach can be seen as a sort of bootstrapping and aggregated neural network model, i.e., bagging, which allows the neural network model to capture distinct letter shapes across various languages.

1(a)

As we do not have a specific learning rate requirement, we set the initial learning rate to 0.001 based on the preliminary experiments and use the exponential decay strategy to plan our learning rate. Then, consider following different filter numbers.

| lr initial | filter | val loss | val acc | test loss | test acc |
|------------|------------------------------|----------|---------|-----------|----------|
| 0.001 | c(8, 16) | 0.6914 | 0.5778 | 0.6914 | 0.5774 |
| 0.001 | c(16, 32) | 0.5957 | 0.7242 | 0.5840 | 0.7480 |
| 0.001 | c(8, 16, 32) | 0.3534 | 0.8980 | 0.3405 | 0.9134 |
| 0.001 | c(16, 32, 64) | 0.6455 | 0.5784 | 0.6467 | 0.5774 |
| 0.001 | c(8, 16, 32, 64) | 0.6771 | 0.5778 | 0.6771 | 0.5774 |
| 0.001 | c(16, 32, 64, 128) | 0.3696 | 0.9032 | 0.3588 | 0.9081 |
| 0.001 | c(16, 32, 64, 128, 256) | 0.3342 | 0.9265 | 0.3100 | 0.9423 |
| 0.001 | c(32, 64, 128, 256, 512) | 0.5700 | 0.5778 | 0.5663 | 0.5774 |
| 0.001 | c(16, 32, 64, 128, 256, 512) | 0.3927 | 0.8653 | 0.3877 | 0.8556 |

From the above results, it is observed that filter sequences c(16, 32, 64, 128, 256) exhibit superior performance with an accuracy of 0.9423. Moving on to 1(b), we also consider filter sequences c(8, 16, 32), c(16, 32, 64, 128), and c(16, 32, 64, 128, 256, 512) as they demonstrate commendable accuracy. The following loss vs accuracy plot corresponds to the model with filter sequences c(16, 32, 64, 128, 256). It does not show an over-fitting trend, and the learning rate strategy works properly.



1(b)

| lr_initial | filter | rotation | val_loss | val_acc | test_loss | test_acc |
|------------|------------------------------|----------|----------|---------|-----------|----------|
| 0.001 | c(8, 16, 32) | TRUE | 0.6605 | 0.5778 | 0.6628 | 0.5774 |
| 0.001 | c(16, 32, 64, 128) | TRUE | 0.6742 | 0.5778 | 0.6746 | 0.5774 |
| 0.001 | c(16, 32, 64, 128, 256) | TRUE | 0.6745 | 0.5778 | 0.6746 | 0.5774 |
| 0.001 | c(16, 32, 64, 128, 256, 512) | TRUE | 0.5519 | 0.7860 | 0.5365 | 0.7822 |

| lr_initial | filter | flipping | loss | val_acc | test_loss | test_acc |
|------------|------------------------------|----------|--------|---------|-----------|----------|
| 0.001 | c(8, 16, 32) | TRUE | 0.6916 | 0.5778 | 0.6916 | 0.5774 |
| 0.001 | c(16, 32, 64, 128) | TRUE | 0.6352 | 0.5778 | 0.6334 | 0.5774 |
| 0.001 | c(16, 32, 64, 128, 256) | TRUE | 0.6622 | 0.5778 | 0.6630 | 0.5774 |
| 0.001 | c(16, 32, 64, 128, 256, 512) | TRUE | 0.6746 | 0.5778 | 0.6754 | 0.5774 |

Rotation and flipping do not improve the performance.

1(c)

| lr_initial | filter | | val_loss | val_acc | test_loss | test_acc |
|------------|------------------------------|------------|----------|---------|-----------|----------|
| 0.001 | c(16, 32, 64, 128, 256) | batch_norm | 0.0948 | 0.9889 | 0.0868 | 0.9948 |
| 0.001 | c(16, 32, 64, 128, 256, 512) | batch_norm | 0.0497 | 0.9918 | 0.0421 | 0.9948 |
| 0.001 | c(16, 32, 64, 128, 256) | sep_conv | 0.6810 | 0.5778 | 0.6811 | 0.5774 |
| 0.001 | c(16, 32, 64, 128, 256, 512) | sep_conv | 0.6810 | 0.5778 | 0.6811 | 0.5774 |
| 0.001 | c(16, 32, 64, 128, 256) | resid_con | 0.6693 | 0.5778 | 0.6692 | 0.5774 |
| 0.001 | c(16, 32, 64, 128, 256, 512) | resid_con | 0.6616 | 0.5778 | 0.6592 | 0.5774 |

Batch normalization significantly improves the model's performance. Following that, we identified the best model with the filter sequence c(16, 32, 64, 128, 256) and incorporated batch normalization into parts d and e.


1(d)

From the above result, choose a model with filter c(16,32), plus rotation to provide a visualization of the interesting activation layers. The following figure is a summary of this model.

Model: "model"

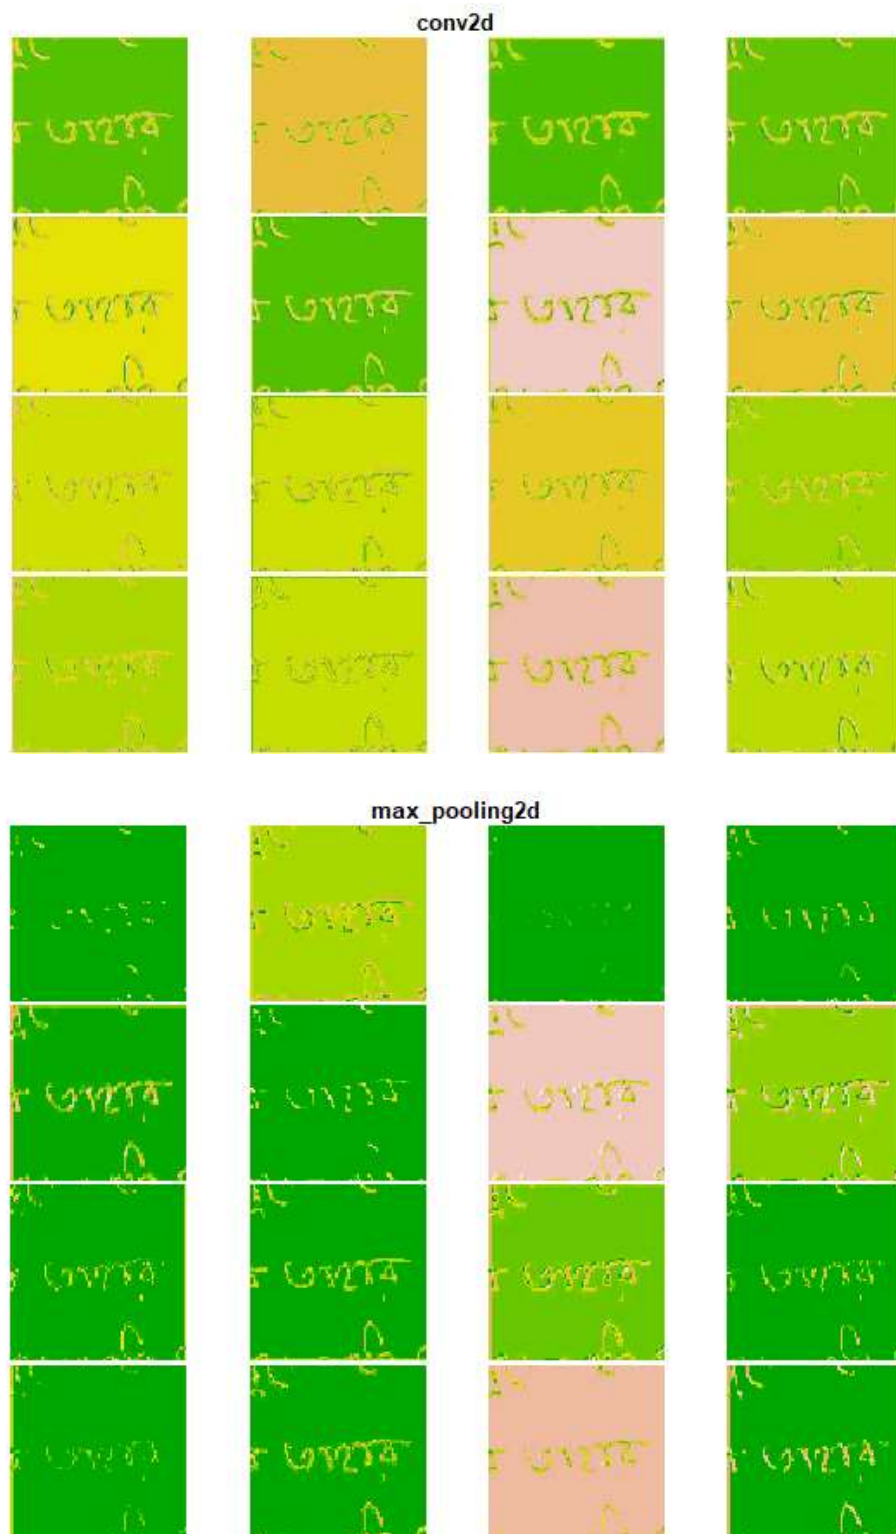
| Layer (type) | Output Shape | Param # | Trainable |
|--|-----------------------|---------|-----------|
| input_1 (InputLayer) | [(None, 150, 150, 1)] | 0 | Y |
| conv2d (Conv2D) | (None, 150, 150, 16) | 144 | Y |
| batch_normalization (BatchNormalization) | (None, 150, 150, 16) | 64 | Y |
| re_lu (ReLU) | (None, 150, 150, 16) | 0 | Y |
| max_pooling2d (MaxPooling2D) | (None, 75, 75, 16) | 0 | Y |
| conv2d_1 (Conv2D) | (None, 75, 75, 32) | 4608 | Y |
| batch_normalization_1 (BatchNormalization) | (None, 75, 75, 32) | 128 | Y |
| re_lu_1 (ReLU) | (None, 75, 75, 32) | 0 | Y |
| max_pooling2d_1 (MaxPooling2D) | (None, 38, 38, 32) | 0 | Y |
| conv2d_2 (Conv2D) | (None, 38, 38, 64) | 18432 | Y |
| batch_normalization_2 (BatchNormalization) | (None, 38, 38, 64) | 256 | Y |
| re_lu_2 (ReLU) | (None, 38, 38, 64) | 0 | Y |
| max_pooling2d_2 (MaxPooling2D) | (None, 19, 19, 64) | 0 | Y |
| conv2d_3 (Conv2D) | (None, 19, 19, 128) | 73728 | Y |
| batch_normalization_3 (BatchNormalization) | (None, 19, 19, 128) | 512 | Y |
| re_lu_3 (ReLU) | (None, 19, 19, 128) | 0 | Y |
| max_pooling2d_3 (MaxPooling2D) | (None, 10, 10, 128) | 0 | Y |
| conv2d_4 (Conv2D) | (None, 10, 10, 256) | 294912 | Y |
| batch_normalization_4 (BatchNormalization) | (None, 10, 10, 256) | 1024 | Y |
| re_lu_4 (ReLU) | (None, 10, 10, 256) | 0 | Y |
| max_pooling2d_4 (MaxPooling2D) | (None, 5, 5, 256) | 0 | Y |
| flatten (Flatten) | (None, 6400) | 0 | Y |
| dense (Dense) | (None, 32) | 204800 | Y |
| batch_normalization_5 (BatchNormalization) | (None, 32) | 128 | Y |
| re_lu_5 (ReLU) | (None, 32) | 0 | Y |
| dense_1 (Dense) | (None, 1) | 33 | Y |
| Total params: 598769 (2.28 MB) | | | |
| Trainable params: 597713 (2.28 MB) | | | |
| Non-trainable params: 1056 (4.12 KB) | | | |

Displaying the test picture:




Visualizing every channel in every intermediate activation:

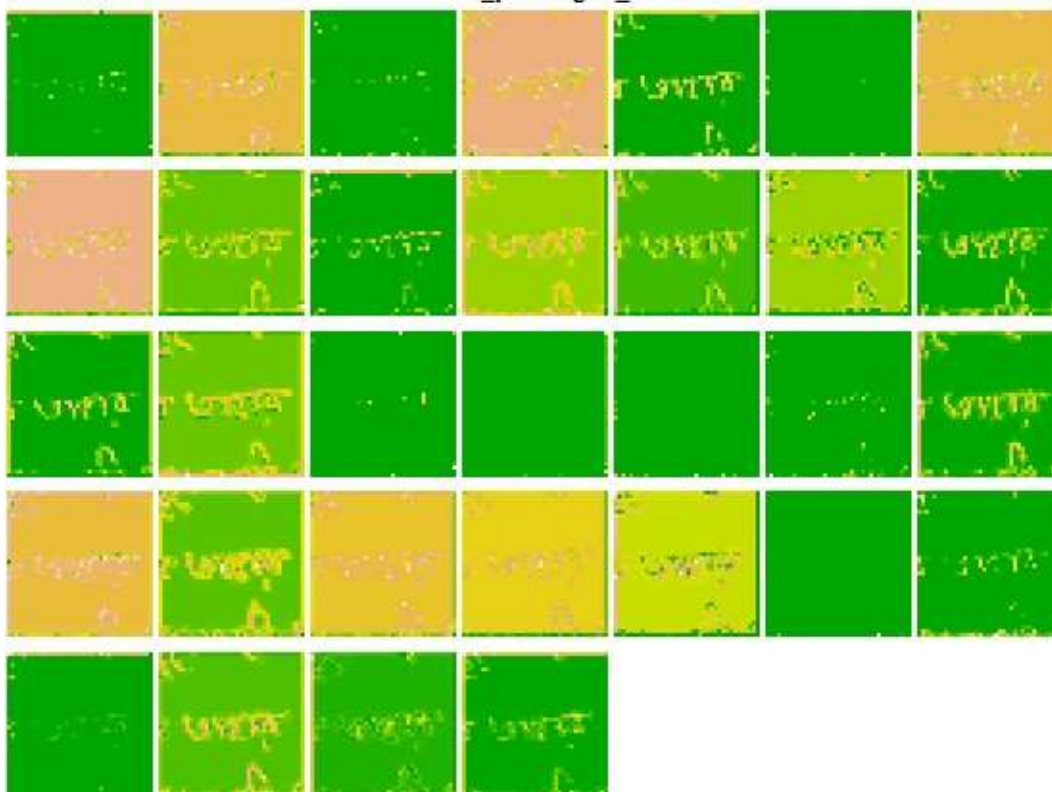
the features extracted by a layer become increasingly abstract with the depth of the layer. The activations of higher layers carry less and less information about the specific input being seen and more and more information about the target, which implies that our model is good.




conv2d_1



max_pooling2d_1



conv2d_4




max_pooling2d_4



1(e)

Visualizing the Probability of Activations:



Visualizing Activations with Context:



Q2

2(a)

As we do not have a specific learning rate requirement, we set the initial learning rate to 0.001 based on the preliminary experiments and use the exponential decay strategy to plan our learning rate. Then we consider following different filter numbers.

| lr initial | filter | val loss | val acc | test loss | test acc |
|------------|----------------------------|----------|---------|-----------|----------|
| 0.001 | c(8, 16) | 0.3078 | 0.8929 | 0.3195 | 0.8750 |
| 0.001 | c(16, 32) | 0.3207 | 0.9127 | 0.3075 | 0.9286 |
| 0.001 | c(8, 16, 32) | 0.3532 | 0.8770 | 0.3890 | 0.8214 |
| 0.001 | c(16, 32, 64) | 0.5048 | 0.8135 | 0.5017 | 0.8750 |
| 0.001 | c(8, 16, 32, 64) | 0.3330 | 0.8611 | 0.4089 | 0.8036 |
| 0.001 | c(16, 32, 64, 128) | 0.1675 | 0.9365 | 0.2483 | 0.9464 |
| 0.001 | c(32, 64, 128, 256) | 0.2123 | 0.9167 | 0.2645 | 0.9464 |
| 0.001 | c(8, 16, 32, 64, 128) | 0.2863 | 0.8929 | 0.3229 | 0.9107 |
| 0.001 | c(8, 16, 32, 64, 128, 256) | 0.2656 | 0.8929 | 0.3304 | 0.8929 |

From the above results, it is observed that filter sequences c(16, 32, 64, 128) and c(32, 64, 128, 256) exhibit superior performance with an accuracy of 0.9464. Moving on to 2(b), we also consider filter sequences c(16, 32) and c(8, 16, 32, 64, 128) as they demonstrate commendable accuracy.

2(b)

| lr initial | filter | rotation | val loss | val acc | test loss | test acc |
|------------|-----------------------|----------|----------|---------|-----------|----------|
| 0.001 | c(16, 32) | T | 0.2531 | 0.9286 | 0.2007 | 0.9643 |
| 0.001 | c(16, 32, 64, 128) | T | 0.3518 | 0.8810 | 0.2991 | 0.9107 |
| 0.001 | c(32, 64, 128, 256) | T | 0.4650 | 0.8571 | 0.4813 | 0.8214 |
| 0.001 | c(8, 16, 32, 64, 128) | T | 0.2729 | 0.9286 | 0.2159 | 0.9643 |

Here, factor = 0.2 for the rotation parameter. And rotation improves the performance of the models with filter sequence c(16, 32) and c(8, 16, 32, 64, 128).

| lr initial | filter | flipping | val loss | val acc | test loss | test acc |
|------------|-----------------------|----------|----------|---------|-----------|----------|
| 0.001 | c(16, 32) | T | 0.3506 | 0.8452 | 0.5536 | 0.7143 |
| 0.001 | c(16, 32, 64, 128) | T | 0.5232 | 0.7381 | 0.5771 | 0.6786 |
| 0.001 | c(32, 64, 128, 256) | T | 0.2013 | 0.9444 | 0.3463 | 0.8214 |
| 0.001 | c(8, 16, 32, 64, 128) | T | 0.3652 | 0.8730 | 0.4586 | 0.8214 |

Flipping does not improve the performance.

2(c)

| lr initial | filter | | val loss | val acc | test loss | test acc |
|------------|-----------------------|------------|----------|---------|-----------|----------|
| 0.001 | c(16, 32) | batch_norm | 1.0796 | 0.4960 | 0.6892 | 0.5357 |
| 0.001 | c(16, 32, 64, 128) | batch_norm | 0.8139 | 0.5040 | 0.6911 | 0.4643 |
| 0.001 | c(32, 64, 128, 256) | batch_norm | 0.6849 | 0.6508 | 0.6694 | 0.3571 |
| 0.001 | c(8, 16, 32, 64, 128) | batch_norm | 1.1362 | 0.5040 | 0.7172 | 0.4643 |
| 0.001 | c(16, 32) | sep_conv | 0.3154 | 0.9444 | 0.2650 | 0.9464 |
| 0.001 | c(16, 32, 64, 128) | sep_conv | 0.6948 | 0.5040 | 0.6965 | 0.4643 |
| 0.001 | c(32, 64, 128, 256) | sep_conv | 0.6939 | 0.5040 | 0.6937 | 0.4643 |
| 0.001 | c(8, 16, 32, 64, 128) | sep_conv | 0.6931 | 0.5040 | 0.6938 | 0.4643 |
| 0.001 | c(16, 32) | resid_conn | 0.6021 | 0.7302 | 0.5768 | 0.7143 |
| 0.001 | c(16, 32, 64, 128) | resid_conn | 0.4415 | 0.8056 | 0.3955 | 0.8571 |
| 0.001 | c(32, 64, 128, 256) | resid_conn | 0.1950 | 0.9524 | 0.1437 | 0.9643 |
| 0.001 | c(8, 16, 32, 64, 128) | resid_conn | 0.4551 | 0.8294 | 0.4196 | 0.8571 |

The skip layers and residual connections, batch/layer normalization, and separable convolutions do not improve the performance in this situation.

2(d)


From the above result, choose a model with filter c(16,32), plus rotation to Provide a visualization of the interesting activation layers. The following figure is a summary of this model.

Model: "model"

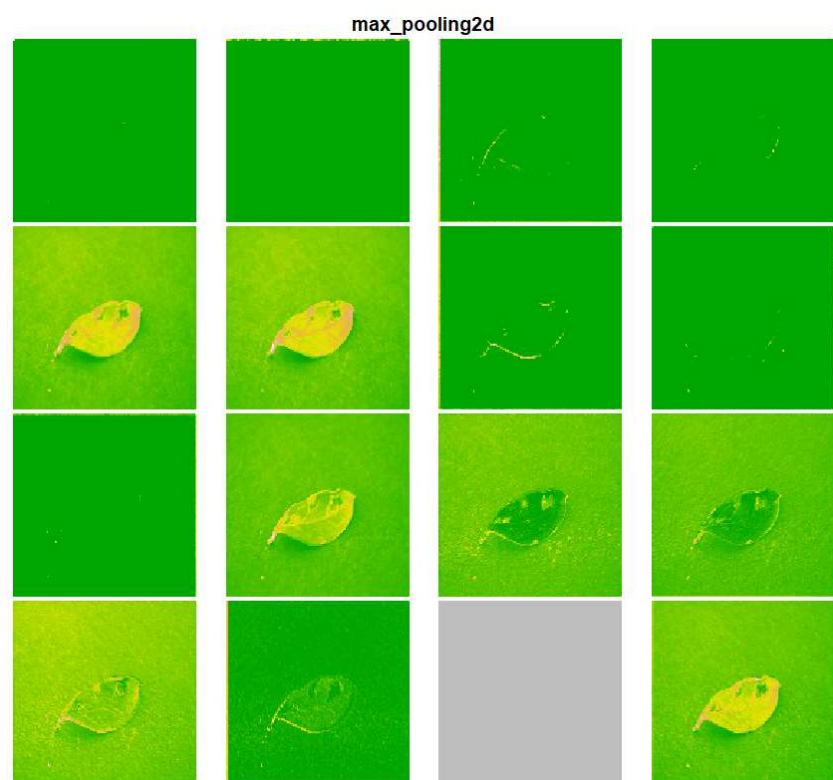
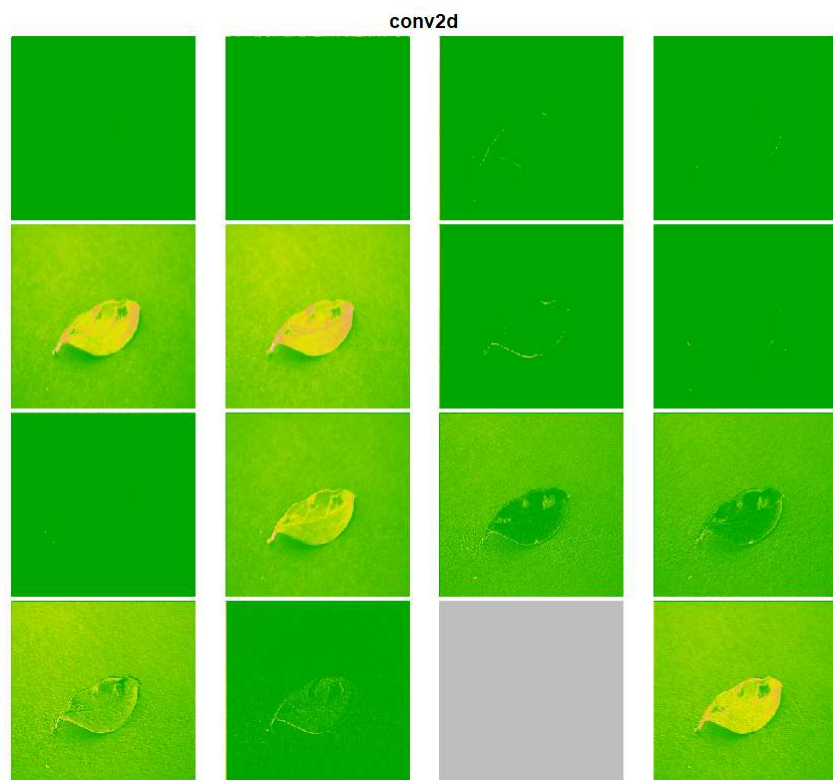
| Layer (type) | Output Shape | Param # | Trainable |
|----------------------------------|-----------------------|---------|-----------|
| input_1 (InputLayer) | [(None, 256, 256, 3)] | 0 | Y |
| random_rotation (RandomRotation) | (None, 256, 256, 3) | 0 | Y |
| conv2d (Conv2D) | (None, 256, 256, 16) | 448 | Y |
| max_pooling2d (MaxPooling2D) | (None, 128, 128, 16) | 0 | Y |
| conv2d_1 (Conv2D) | (None, 128, 128, 32) | 4640 | Y |
| max_pooling2d_1 (MaxPooling2D) | (None, 64, 64, 32) | 0 | Y |
| flatten (Flatten) | (None, 131072) | 0 | Y |
| dense (Dense) | (None, 32) | 4194336 | Y |
| dense_1 (Dense) | (None, 1) | 33 | Y |

=====
Total params: 4199457 (16.02 MB)
Trainable params: 4199457 (16.02 MB)
Non-trainable params: 0 (0.00 Byte)

Displaying the test picture:



Visualizing every channel in every intermediate activation:



conv2d_1




max_pooling2d_1



2(e)

Visualizing the Probability of Activations:



Visualizing Activations with Context:




Here, it's interesting to note that the black spots on the diseased tree leaves are strongly activated. this is probably how the network can tell the difference between disease and no disease.

Q3

3(a)

Model Overview

- We utilized U-Net for the segmentation of neither (background), lung, or airway.
- Input: 128 x 128 x 1 input images (rescaled and greyscale).
 - Trained on 26*200 images (sampling 200 slices from each CT)
- Model: U-Net Architecture



- Settings:
 - Adam optimizer, learning rate = 0.001
 - Categorical cross-entropy with 3 categories (neither/lung/airway)
 - Accuracy metric
 - 10 epochs, 32 batch size

Model Training Loss/Accuracy

- Loss: 0.0345
- Accuracy: 0.9877

Model Validation Loss/Accuracy

- Loss: 0.0385
- Accuracy: 0.9862

Model Test Loss/Accuracy (20% split)

- Loss: 0.0369
- Accuracy: 0.9866

3(b)

Image 27 Loss/Accuracy

- Loss: 0.0495
- Accuracy: 0.9817

Confusion Matrix


| | | Ground Truth | | |
|------------|---------|--------------|--------|--------|
| | | Neither | Lung | Airway |
| | Neither | 2830427 | 46265 | 562 |
| Prediction | Lung | 10528 | 383138 | 3 |
| | Airway | 1217 | 1469 | 3191 |

Slices of the images:


From left to right: x(data), neither(truth), lung(truth), airway(truth)



From left to right: neither(prediction), lung(prediction), airway(prediction)



Another image:



The U-Net model provided good results: high accuracy across training, validation, and test datasets. However, because of the imbalanceness of our data, classifying them as neither or lung achieves a high accuracy score. From the confusion matrix, we noticed that the prediction accuracy for the airway is not consistent with the overall 98% accuracy.

To address this issue and for future investigations, it would be nice to use alternative metrics, such as the multiclass-F1 score, and also put weights to minor categories (airway in our case) to provide more accurate results.